

Machine Learning

Pawel Wocjan

University of Central Florida


Fall 2020

Linear Regression

- ▶ Consider the following toy example.
- ▶ It has been known that crickets (an insect species) chirp more frequently on hotter days than on cooler days.
- ▶ Professional and amateur scientists have cataloged data on chirps-per-minute and temperature.
- ▶ Using this data, you want to explore this relationship.


Linear Regression

- First, examine your data by plotting it:



Linear Regression

- ▶ You could draw a single straight line like the following to approximate this relationship between chirps and temperature.



Linear Regression

- ▶ The line doesn't pass through every dot, but the line does clearly show the relationship between chirps and temperature.
- ▶ Using the equation for a line, you could write down this relationship as follows:

$$y = mx + b$$

where:

- ▶ y is the temperature in Celsius – the value we're trying to predict.
- ▶ m is the slope of the line.
- ▶ x is the number of chirps per minute – the value of our input feature.
- ▶ b is the y -intercept.

Linear Regression

- ▶ By convention in ML, you'll write the equation for a model slightly differently:

$$\hat{y} = b + w_1x_1$$

where:

- ▶ y is the predicted label (a desired output).
- ▶ b is the bias (the y -intercept), sometimes referred to as w_0 .
- ▶ w_1 is the weight of feature 1. Weight is the same concept as the “slope” m in the traditional equation of a line.
- ▶ x_1 is a feature (a known input).

Linear Regression

- ▶ To **infer** (predict) the temperature y' for a new chirps-per-minute value x_1 , just substitute the x_1 value into this model.
- ▶ Although this model uses only one feature, a more sophisticated model might rely on multiple features, each having a separate weight w_1 , w_2 , etc.
- ▶ For example, a model that relies on three features might look as follows:

$$\hat{y} = b + w_1x_1 + w_2x_2 + w_3x_3$$

Key Terms

- ▶ bias
- ▶ inference
- ▶ linear regression
- ▶ weight