

Gradient for linear regression

Pawel Wocjan

February 3, 2020

Abstract

We compute the gradient for linear regression.

1 Linear regression with single feature

Let $w \in \mathbb{R}$ and $b \in \mathbb{R}$ be the weight and bias for linear regression. Given $x \in \mathbb{R}$, the predicted value is

$$\hat{y} = wx + b. \quad (1)$$

Assume that the correct value for x is $y \in \mathbb{R}$. Then the squared error loss is given by

$$\mathcal{L} = \frac{1}{2}(\hat{y} - y)^2. \quad (2)$$

The gradient of the loss function is

$$\nabla \mathcal{L} = \begin{pmatrix} \frac{\partial \mathcal{L}}{\partial b} \\ \frac{\partial \mathcal{L}}{\partial w} \end{pmatrix} \in \mathbb{R}^2. \quad (3)$$

Using the chain rule, we compute the bias component of the gradient

$$\frac{\partial \mathcal{L}}{\partial w} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial b} = (\hat{y} - y) \cdot 1 \quad (4)$$

and the weight component of the gradient

$$\frac{\partial \mathcal{L}}{\partial w} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial w} = (\hat{y} - y) \cdot x. \quad (5)$$

Putting these together, we obtain

$$\nabla \mathcal{L} = \begin{pmatrix} 1 \\ x \end{pmatrix} \cdot (\hat{y} - y) \quad (6)$$

2 Linear regression for multiple features

2.1 Single example

Having dealt with the simple case of linear regression with a single feature in the previous section, we now consider the general case of linear regression with n features.

Let $w = (w_1, \dots, w_n)^T \in \mathbb{R}^n$ and $b \in \mathbb{R}$ be the weight (column) vector and the bias term, respectively. Given the feature (column) vector $x = (x_1, \dots, x_n)^T \in \mathbb{R}^n$, the predicted value by the linear regression model is

$$\hat{y} = \sum_{j=1}^n w_j x_j + b. \quad (7)$$

It will be convenient to treat the weights w_j and the bias b in a unified way. To do this, set $w = (w_0, w_1, \dots, w_n)^T \in \mathbb{R}^{n+1}$ with $w_0 = b$ and $x = (1, x_1, \dots, x_n)^T \in \mathbb{R}^{n+1}$. The predicted value is then given by

$$\hat{y} = \sum_{j=0}^n w_j x_j. \quad (8)$$

Note that the prediction \hat{y} can also be expressed as the dot product $x^T w$.

The loss is equal to

$$\mathcal{L} = \frac{1}{2}(\hat{y} - y)^2 = \frac{1}{2}(x^T w - y)^2. \quad (9)$$

Using the chain rule to compute the partial derivatives $\partial \mathcal{L} / \partial w_j$ for $j = 0, \dots, n$, we obtain the expression for the gradient

$$\nabla \mathcal{L} = x(\hat{y} - y) = x(x^T w - y) \in \mathbb{R}^{n+1}. \quad (10)$$

2.2 Multiple examples

Let $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)}) \in \mathbb{R}^{n+1} \times \mathbb{R}$ be the training example in a batch. The loss for the i th example is

$$\mathcal{L}^{(i)} = \frac{1}{2}(\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{2}(x^{(i)T} w - y^{(i)})^2 \quad (11)$$

and the gradient of the loss for the i th example is

$$\nabla \mathcal{L}^{(i)} = x^{(i)}(\hat{y}^{(i)} - y^{(i)}) = x^{(i)}(x^{(i)T} w - y^{(i)}). \quad (12)$$

The mean squared error loss for the batch is

$$\text{MSE} = \frac{1}{m} \sum_{i=1}^m \mathcal{L}^{(i)} \quad (13)$$

and the gradient of the MSE for the batch is

$$\nabla \text{MSE} = \frac{1}{m} \sum_{i=1}^m \nabla \mathcal{L}^{(i)} \quad (14)$$

$$= \frac{1}{m} \sum_{i=1}^m x^{(i)} (x^{(i)T} w - y^{(i)}). \quad (15)$$

The summation corresponds to a for-loop. To write vectorized code making it possible to process the examples in the batch in parallel, we have to avoid the explicit summation. To this end, define the matrix

$$X = (x^{(1)} \dots x^{(m)}) \in \mathbb{R}^{(n+1) \times m} \quad (16)$$

whose columns are the feature (column) vectors $x^{(i)} \in \mathbb{R}^{n+1}$ and the vector

$$y = (y^{(1)} \dots y^{(m)}) \in \mathbb{R}^{1 \times m} \quad (17)$$

whose entries are the labels $y^{(i)} \in \mathbb{R}$.

It is a little bit tricky, but it can be shown that

$$\nabla \text{MSE} = X(X^T w - y^T). \quad (18)$$

This provides the basis for the vectorized implementation of mini-batch gradient descent in the notebook

<https://colab.research.google.com/drive/1qBxfTPoNcSFvpwu1ND11V6cHEqL3aQ1->
using the command `numpy.dot`

<https://docs.scipy.org/doc/numpy/reference/generated/numpy.dot.html>.

(Note that in the code we use X^T instead of X and y^T instead of y . This is because the first dimension is the batch dimension. I will explain this in class.)